

EP24040 - Computational Physics

Kirit Makwana
Department of Physics, IIT Hyderabad

Neural networks

Motivation

- There are a whole slew of numerical methods that aim to approximate the exact solution of equations with some discretization procedure
- In the case of Monte-Carlo techniques, there is a randomization procedure
- The approximation can also be done using spectral methods - methods in which the problem is transformed from real space to another more suitable space
- Fourier transforms, Hermite transform, Chebyshev transforms, etc
- Further still, Monte Carlo methods have also been modified to make computations faster using greedy algorithms
- All these techniques are extensively used in Physics, as well as in other fields

Classification

- A new class of problem has arisen due to large availability of data
- The problem is classification of data
- Often times we have lots of data and need to organize it into human relevant classes
- For example, identifying different animals from image data, or converting a signal into an alphabet (speech recognition)
- Classifying the types of galaxies (spiral, elliptical)



Spiral

Elliptical

Irregular

Classification

- Analyzing noisy data, removing or reducing noise, predicting disruptive events
- In these problems we have a lot of input data, and from past experience we know what output we get in the case of input data
- We need to make a model such that given some input data, we can predict the outcome
- For example, sample image data and identify the animal in that



Input output vectors

- Any form of input can be converted into a vector
- Pixels from a photo can be arranged in an input vector $x \in \mathbb{R}^n$
- The output can be another vector $y \in \mathbb{R}^m$
- The output can be very simple such as a two element vector, signifying whether the animal is cat or dog

$$\begin{bmatrix} 1 \\ 0 \end{bmatrix} \Rightarrow \text{Cat or } \begin{bmatrix} 0 \\ 1 \end{bmatrix} \Rightarrow \text{Dog} \quad (1)$$

Linear transformation

- A linear model takes a linear transformation from $x \in \mathbb{R}^n$ to $y \in \mathbb{R}^m$
- This can be done by simply matrix multiplication

$$y = Ax \quad (2)$$

- where A is an $m \times n$ matrix which we are interested in finding
- Lets say we have 3 input vector $x_1 = [1, -1]$, $x_2 = [0, 1]$, and $x_3 = [1, 1]$
- Lets say for these cases, we have the output $y_1 = [0, 1]$, $y_2 = [0, 1]$, $y_3 = [1, 0]$ respectively
- We can try to find an A which will be a 2×2 matrix such that

$$\begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} 1 \\ -1 \end{bmatrix} = \begin{bmatrix} 0 \\ 1 \end{bmatrix} \text{ and } \begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} 0 \\ 1 \end{bmatrix} = \begin{bmatrix} 0 \\ 1 \end{bmatrix} \quad (3)$$

Linear model

- We can see that A is

$$A = \begin{bmatrix} 0 & 0 \\ 2 & 1 \end{bmatrix} \quad (4)$$

- This matrix is uniquely determined by the first two conditions
- But then it does not satisfy the third condition
- This means a linear model will not work in this case
- In general, as the number of cases increases, you will get more conditions, while the matrix will have maximum rank of n
- So the system will be overdetermined and we cannot fit the linear model exactly
- One way to overcome this drawback is to go to a nonlinear model

Nonlinear model

- Take the output of a linear model and apply a nonlinear function to it

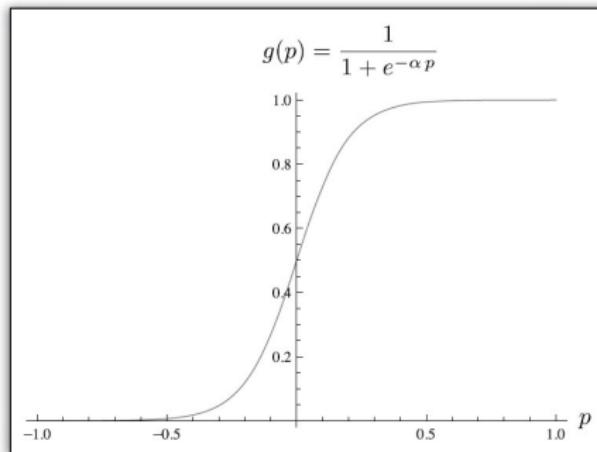
$$\mathbf{z} = g(\mathbf{y}) = \begin{bmatrix} g(y_1) \\ g(y_2) \\ \vdots \\ g(y_m) \end{bmatrix} \quad (5)$$

- Defining $g(p)$ gives extra flexibility to fit all the required conditions, ex. the third condition above
- The input and outputs can be standardized for ease of operation
- Set $g(p)$ to lie between 0 and 1

Sigmoid function

- The sigmoid function is a continuous version of the step function

$$g(p) = \frac{1}{1 + e^{-\alpha p}} \quad (6)$$



- α sets the width of the sigmoid function, higher α makes the function sharper

Setting α

- The elements of the nonlinear output \mathbf{z} are

$$z_i = g((A\mathbf{x})_i) = g\left[\sum_{j=1}^n A_{ij}x_j\right] \quad (7)$$

- Now the argument of function g should span the range of the width of the sigmoid function, so that the output g takes the full range of 0 to 1
- Assume that the elements of matrix A will be of the order unity
- Then the maximum element of $A\mathbf{x}$ will be approximately $n\max(x_i)$
- Then take α as

$$\alpha = \frac{10}{n\max(x_i)} \quad (8)$$

- Also, we can shift the input numbers $x_i \rightarrow x_i - \text{med}(x_i)$ such that the maximum and minimum are equal and opposite in signs
- This way we can handle input with any magnitude \mathbf{x}

Training

- Now the next step is to find the matrix A
- This process is called as training. In this we provide training examples to the model
- These examples are basically pairs of $[x^k, y^k]_{k=1}^T$ where x^k is the k-th input vector and y^k is its desired output vector
- There are T training examples, labelled by k , and our goal is to set A such that we get y^k for given x^k
- For this we define an error function f

$$f^k = ||g(Ax^k) - y^k||^2 \text{ for each case } j = 1, 2, \dots, T \quad (9)$$

- The norm of a vector can be defined as the square root of the sum of the squares of its elements, i.e.,

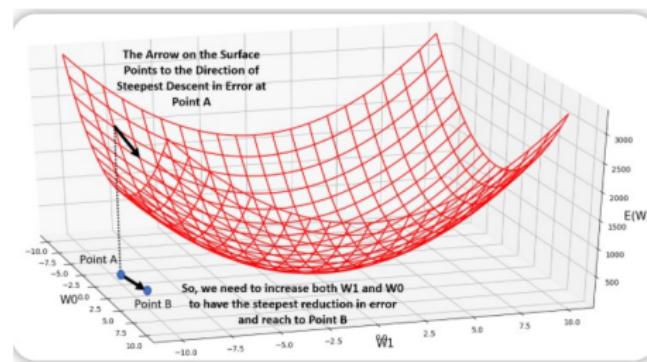
$$f^k = \sum_{i=1}^m \left[g\left(\sum_{s=1}^n A_{is} x_s^k\right) - y_i^k \right]^2 \quad (10)$$

Training

- This function f can be thought of as a function of the matrix A_{ij} because that is the only variable, rest all quantities come from the training set
- Hence, the objective then is to adjust A_{ij} such that we can minimize f for each training example
- f is a positive definite quantity, and this problem is of minimization of f
- A well-known method is known as the steepest descent (gradient descent) method

Gradient descent

- Consider a function of 2 variables, this can be represented as a surface
- This function contains has a minima



- If we have a function $f(x)$ of which we want to find minimum, then starting from some position $x_0 = (x_0, y_0)$

$$x_{n+1} = x_n - \eta \nabla f \quad (11)$$

- η is the step size, also called as “rate of learning”

Steepest descent

- Ex. find the minimum of $f(x, y) = x^2 + 4y^2 - 4x - 24y + 40$
- This actually comes from $(x - 2)^2 + 4(y - 3)^2$, so the minimum is at $(x, y) = (2, 3)$
- Learning rate can be faster with higher η
- However, too high η might lead to method not converging or instability
- Also called steepest descent because change is maximum in direction of gradient
- Convergence is determined when the steps become smaller than some threshold

Training model

- Going back to the case of finding the matrix A in our model
- The function to be minimized is

$$f(A_{ij}) = \sum_{i=1}^m \left[g\left(\sum_{j=1}^n A_{ij}x_j\right) - y_i \right]^2 \quad (12)$$

- given any training pair $(\mathbf{x}^k, \mathbf{y}^k)$ - suppressing the superscript k for simplicity
- Define $b_i = \sum_{j=1}^n A_{ij}x_j$ and $z_i = g(b_i)$, i goes from 1 to m
- Now we want to find the values of matrix elements A_{pq} which minimizes the error f
- For using gradient descent we need

$$\frac{\partial f}{\partial A_{pq}} = \sum_{i=1}^m 2(z_i - y_i) \frac{\partial z_i}{\partial A_{pq}} \quad (13)$$

Training model

- We have

$$\frac{\partial z_i}{\partial A_{pq}} = g'(b_i) \frac{\partial b_i}{\partial A_{pq}} \quad (14)$$

$$\frac{\partial b_i}{\partial A_{pq}} = \frac{\partial}{\partial A_{pq}} \sum_{j=1}^m A_{ij} x_j \quad (15)$$

$$= \sum_{j=1}^m \delta_{ip} \delta_{jq} x_j \quad (16)$$

$$= \delta_{ip} x_q \quad (17)$$

- Also can be shown that

$$g'(p) = \alpha g(p)(1 - g(p)) \quad (18)$$

Training model

- Putting this together we get

$$\frac{\partial f}{\partial A_{pq}} = \sum_{i=1}^m 2(z_i - y_i) \alpha g(b_i)(1 - g(b_i)) \delta_{ip} x_q \quad (19)$$

- which gives

$$\frac{\partial f}{\partial A_{pq}} = 2(z_p - y_p) \alpha z_p (1 - z_p) x_q \quad (20)$$

- This makes the calculation of the matrix element easy

Algorithm

- Initialize random A of order unity
- Loop over each training pair
- Calculate z using the existing A matrix
- Loop over all elements of A_{pq} and update with

$$A_{pq} \rightarrow A_{pq} - \eta \left[2(z_p - y_p) \alpha z_p (1 - z_p) x_q \right] \quad (21)$$

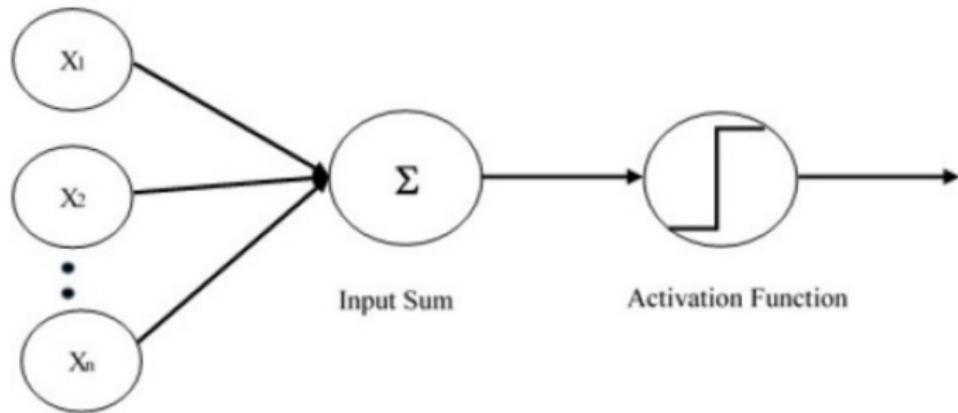
- Taking one gradient descent step over all the training examples is called an “epoch”
- Run the training for several epochs till the error is reduced below some threshold or the accuracy of the model reaches desired level
- The accuracy can be defined as the percentage of predictions it gets right
- After going through all training pairs and epochs, the model is trained and can be used for new data

Training model

- This training is not expected to produce exact results
- The outcome of the training will depend on the order in which the test cases are given
- Ideally after giving the training examples multiple times and in random order, the neural network A should converge
- However, this may not be the case always and it may not give good agreement for all the training examples
- After the model is trained on the training set, it can be tested on a testing set to check how well it performs
- If the performance is good, it can be used on new and untested data

Hidden layers

- If we are given an input vector with 10 elements, i.e., $n = 10$ and the output is just a single element, i.e., $m = 1$ then
- The matrix A will be 10×1 matrix with 10 parameters to adjust
- This can be represented diagrammatically as



Hidden layers

- If there is a lot of data which is making the problem overconstrained, then we require more freedom in the neural network model
- We can add an intermediate layer of size k
- The input $\mathbf{x} \in R^n$ gets multiplied by matrix B which is $k \times n$ giving a $\tilde{\mathbf{z}} \in R^k$

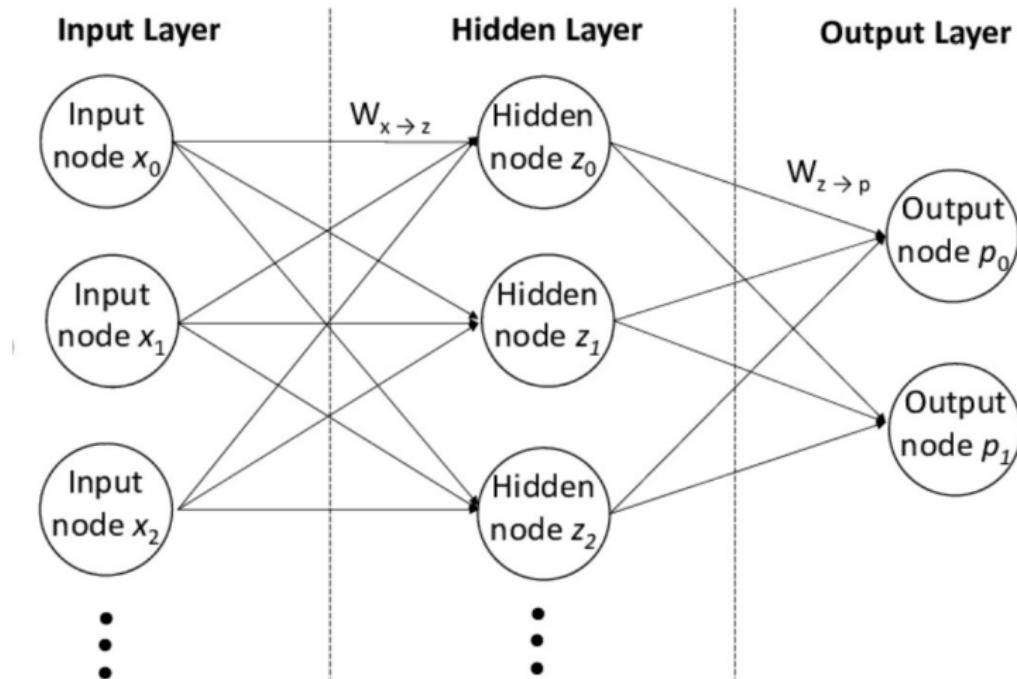
$$\tilde{\mathbf{z}} = g(B\mathbf{x}) - \frac{1}{2} \quad (22)$$

- Since the nonlinear function g is applied here, the $-1/2$ changes the mapping from $0 - 1$ to $-1/2$ to $+1/2$
- Now the matrix A of size $m \times k$ is applied on $\tilde{\mathbf{z}}$ to get the final output

$$\mathbf{z} = \bar{g}(A\tilde{\mathbf{z}}) \quad (23)$$

Hidden layers

- The hidden layers can be diagrammatically represented as



Training

- Now both the B and A matrices have to be found by training the model
- The error is again defined as

$$f(A_{ls}, B_{ij}) = \sum_{l=1}^m (z_l - y_l)^2 \quad (24)$$

- Using the steepest gradient technique as before, we can calculate the partial derivative of this w.r.t. the matrix elements A_{ls} and B_{ij}
- It can be shown that

$$\frac{\partial f}{\partial A_{pq}} = 2\bar{\alpha}(z_p - y_p)z_p(1 - z_p)\tilde{z}_q \quad (25)$$

Training

- Similarly, differentiating w.r.t. B_{pq} it can be shown that

$$\frac{\partial f}{\partial B_{pq}} = \sum_{l=1}^m \sigma_l A_{lp} \alpha \left(\frac{1}{2} + \tilde{z}_p \right) \left(\frac{1}{2} - \tilde{z}_p \right) x_q \quad (26)$$

- where

$$\sigma_l = 2\bar{\alpha}(z_l - y_l)z_l(1 - z_l) \quad (27)$$

- Now we can just modify the update algorithm

$$A_{pq} \rightarrow A_{pq} - \eta \frac{\partial f}{\partial A_{pq}} \quad (28)$$

$$B_{pq} \rightarrow B_{pq} - \eta \frac{\partial f}{\partial B_{pq}} \quad (29)$$

- More generally, the gradient w.r.t. the weights of the neural network nodes are estimated by a method of backward propagation

Different models

- A model with a large number of hidden layers is a deep network
- To generate mode data, noise can be added to the data
- The images can also be transformed, like left-right inversion, or time inversion, etc.
- Convolutional neural network - image data is processed through some convolution in order to account for neighboring/related pixels
- Taking a large hidden layer (high k) versus many hidden layers - mostly based on experience and experimentation
- Steepest descent leads to the local minimum, might not always work or find the best possible network - therefore need to try many different starting points and train multiple times

Examples

- Toy problem - <https://machinelearningmastery.com/develop-your-first-neural-network-with-pytorch-step-by-step/>
- Classify signals into frequency - this can be used for analyzing time dependent data
- Classification of galaxies -
<https://medium.com/analytics-vidhya/multiclass-image-classification-problem-convolutional-neural-networks-93a2a2a2a2a2>
- Identifying and classifying phases of matter based on experimental or simulation results

Example

- Consider the input and output examples

$$[0.45, 0.6, -0.27] \rightarrow 1 \quad (30)$$

$$[0.85, -0.57, 0.14] \rightarrow 0 \quad (31)$$

$$[0.02, 0.96, -0.15] \rightarrow 1 \quad (32)$$

$$[-0.63, 0.79, -0.25] \rightarrow 0 \quad (33)$$

$$[-0.82, 0.15, 0.45] \rightarrow 1 \quad (34)$$

$$[0.48, 0.75, 0.64] \rightarrow 0 \quad (35)$$

(36)

- Train a neural network on this data and try to predict the output for new data

Physics informed neural networks

- Neural networks can also be used to solve differential equations describing physical laws
- For example - suppose we experimentally measure the displacement of a particle as a function of time

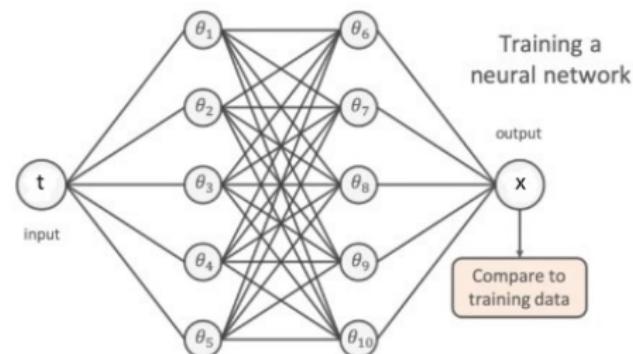
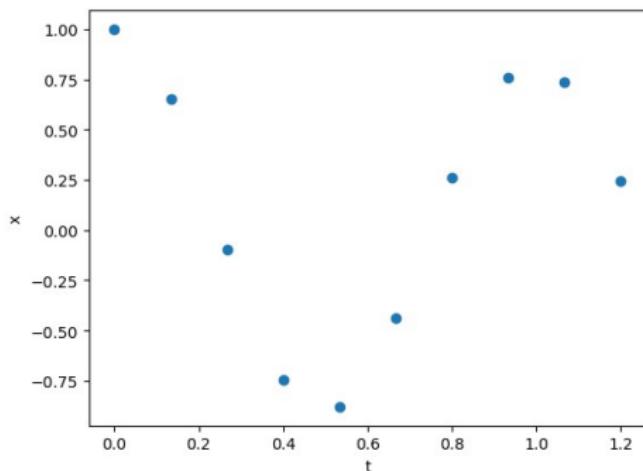


Figure: Ben Mosley

- Now we need to design a neural network that can fit a function $f(t)$ that will approximately give us the value of the displacement $x(t)$

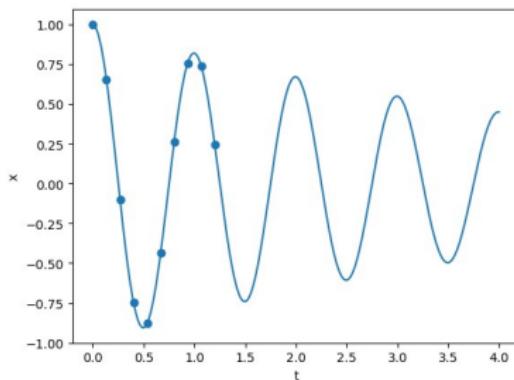
$$f(t) \approx u(t) \quad (37)$$

Physics informed neural networks

- With limited data, the neural network will not be able to extrapolate the displacement to later times
- Rather, if we know that the displacement obeys a differential equation then we can get the full solution

$$m \frac{d^2x}{dt^2} + \mu \frac{dx}{dt} + kx = 0 \quad (38)$$

- This is the equation of a damped harmonic oscillator (SHM if $\mu = 0$) and its solution looks as below



- Consider the damped harmonic oscillator differential equation

$$m \frac{d^2x}{dt^2} + \mu \frac{dx}{dt} + kx = 0 \quad (39)$$

- with the initial condition $x(t = 0) = 1$ and $x'(t = 0) = 0$, i.e. an oscillator starting from the peak displacement
- Setup a neural network NN that takes t as an input and gives x as output
- We have to train the network weight factors θ such that

$$NN(t, \theta) \approx x(t) \quad (40)$$

- We do not need any experimental data to train this network
- The neural network will start generating data itself and the differential equation and boundary conditions will help in training

Loss function

- The network is initialized with random weights
- In each epoch a few input collocation times (t_i , $i=1,2,\dots,N_p$) are chosen randomly over the range of interest (including the boundary point) - it is important to train the model over random time points in order to avoid bias towards early or later times
- The network makes predictions $NN(t_i, \theta)$
- The loss function consists of two parts, the boundary loss term

$$f_b = \lambda_1 \left(NN(t=0, \theta) - 1 \right)^2 + \lambda_2 \left[\frac{dNN}{dt}(t=0, \theta) - 0 \right]^2 \quad (41)$$

- and the differential equation loss is (also called physics-loss)

$$f_p = \frac{\lambda_3}{N_p} \sum_{i=1}^{N_p} \left[m \frac{d^2 NN(t_i)}{dt^2} + \mu \frac{dNN(t_i)}{dt} + kNN(t_i) \right]^2 \quad (42)$$

Grad function

- The weights of the different loss terms $\lambda_1, \lambda_2, \lambda_3$ can be adjusted
- The derivatives dNN/dt and d^2NN/dt^2 need to be calculated

```
from torch.autograd import Variable  
t_collocation = np.random.uniform(0,4, size=(500,1))  
pt_t_collocation = Variable(torch.from_numpy\n(t_collocation).float(), requires_grad=True)
```

- After defining a network $NN(time)$ which takes an input $time$ tensor, it's derivative can be calculated

```
xn=NN(time)  
dxndt=torch.autograd.grad(xn,time)
```

- Similarly, getting the second derivative will involve taking one more grad on the $dxndt$ variable

Wave equation

- In order to solve the wave equation for a quantity $u(x, t)$

$$\frac{\partial^2 u}{\partial x^2} - \frac{1}{c^2} \frac{\partial^2 u}{\partial t^2} = 0 \quad (43)$$

- For this we will define a $NN(x, t)$ and the physics loss will be

$$f_p = \frac{1}{N_p} \sum_i \left[\frac{\partial^2 NN(x_i, t_i)}{\partial x^2} - \frac{1}{c^2} \frac{\partial^2 NN(x_i, t_i)}{\partial t^2} \right] \quad (44)$$

- Where the summation is over all collocation points
- For periodic boundary condition the loss is

$$f_b = \sum_j [NN(x_j, t_j) - NN(x_j + L, t_j)]^2 \quad (45)$$

- where x_j is a boundary points and L is the period

Epilogue

- Computational physics and scientific computing is a relatively young field
- However, its growth is very rapid and it is continuing to develop faster
- Not just numerical techniques but hardware also gets more advanced
- Massively parallel computing, graphical processing units, cloud computing
- Parallelization with OpenMP, MPI, GPU, OpenACC
- Quantum computing
- Other courses
 - ▶ Computational Solid state physics
 - ▶ Computational Particle Physics
 - ▶ Computational Modeling of Biological Systems
 - ▶ Data Science Analysis